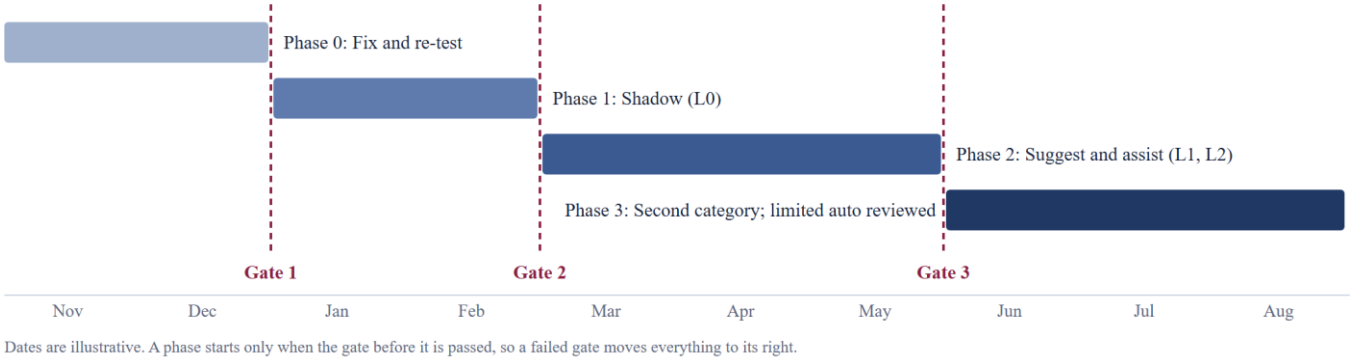
Roadmap with Proof Gates on TradeSentinel

Team013 (Individual) · Jodhdeep Singh · 2025mb26112 · MBA – AI in Business · Management of AI Products · November 2026

This roadmap is about staged learning. Each phase exists to answer a question, and dates are illustrative because a failed gate moves everything after it.



Illustrative timeline. Gate 1 needs Phase 0 to finish.

Phases

Phase	Question it answers	Scope	Autonomy	Main outputs
0. Fix and re-test (Nov to Dec 2026)	Can we close the understated-report gap? Does the design save time?	Synthetic data; fresh test set by a second author; timed trial with five engineers	None	Gate 1 evidence; API route chosen
1. Shadow (Jan to Feb 2027)	Does it work on real tickets, for all client groups?	One flow, tier and venue; masked data; expert labels; incident replay; fairness and injection tests	L0	Gate 2 evidence; real critical share; human baseline
2. Suggest and assist (Mar to May 2027)	Do engineers use it well? Is net payoff positive?	L1 first; L2 only per category after review; execution failures stay L1	L1, then L2	Gate 3 evidence; sign-offs
3. Second category (Jun to Aug 2027)	Is the design reusable? Should anything be automated?	Market data feed alerts, starting at L0; formal review of L3 for lowest-risk categories	L0 for new; L3 reviewed	Scale or assist-only decision

Gate Criteria and Stop Conditions

Gate	Pass requires	If failed
Gate 1	New held-out set: no critical issue below the escalation line; cause in top three at least 70%. Timed trial: under 3 minutes median to first correct triage	Fix and repeat once. If it fails again, return to the design or stop
Gate 2	Category agreement within 5 points of expert-to-expert agreement; no S1 called below S2; no critical issue below the escalation line, including replay; no material fairness gap; masking verified	Stay in shadow. Two consecutive failures: stop or limit to search support
Gate 3	Positive net payoff against current practice; ECE below 0.05; stable override and review time; operations and risk sign-off	Stay at L1 as an assist feature
Reuse decision	Second category live with only small pluggable strategies added	Stay a single-flow feature, not a platform

Ladder Rules

- Levels are set per category, never globally.
- A category drops a level at once on any safety breach; going back up requires passing the gate again.
- Execution failures do not go above L1 in this roadmap.
- L3 covers reversible actions on low-risk categories only.